

You can import your own data into Colab notebooks from your Google Drive account, including from spreadsheets, as well as from Github and many other sources. To learn more about importing data, and how Colab can be used for data science, see the links below under [Working with Data](#).

Business Case: LoanTap Logistic Regression

Problem Statement: Given a set of attributes for an Individual, determine if a credit line should be extended to them. If so, what should the repayment terms be in business recommendations?

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

data = pd.read_csv("https://d2beiqrh929f0.cloudfront.net/public_assets/assets/000/003/549/original/logistic_regression.csv?1651045921")

data.head()
```

	loan_amnt	term	int_rate	installment	grade	sub_grade	emp_title	emp_length	home_ownership	annual_inc	...	open_acc	pub
0	10000.0	36 months	11.44	329.48	B	B4	Marketing	10+ years	RENT	117000.0	...	16.0	
1	8000.0	36 months	11.99	265.68	B	B5	Credit analyst	4 years	MORTGAGE	65000.0	...	17.0	
2	15600.0	36 months	10.49	506.97	B	B3	Statistician	< 1 year	RENT	43057.0	...	13.0	
3	7200.0	36 months	6.49	220.65	A	A2	Client Advocate	6 years	RENT	54000.0	...	6.0	
4	24375.0	60 months	17.27	609.33	C	C5	Destiny Management Inc.	9 years	MORTGAGE	55000.0	...	13.0	

5 rows × 27 columns

```
# features & label
data.columns

Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade', 'sub_grade',
      'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
      'verification_status', 'issue_d', 'loan_status', 'purpose', 'title',
      'dti', 'earliest_cr_line', 'open_acc', 'pub_rec', 'revol_bal',
      'revol_util', 'total_acc', 'initial_list_status', 'application_type',
      'mort_acc', 'pub_rec_bankruptcies', 'address'],
      dtype='object')

# size of data
data.shape

(396030, 27)

# Type of data
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   loan_amnt              396030 non-null float64
1   term                   396030 non-null object
2   int_rate                396030 non-null float64
3   installment             396030 non-null float64
4   grade                   396030 non-null object
5   sub_grade               396030 non-null object
6   emp_title               373103 non-null object
7   emp_length              377729 non-null object
8   home_ownership          396030 non-null object
9   annual_inc              396030 non-null float64
10  verification_status      396030 non-null object
11  issue_d                 396030 non-null object
```

```

12  loan_status      396030 non-null object
13  purpose          396030 non-null object
14  title            394275 non-null object
15  dti              396030 non-null float64
16  earliest_cr_line 396030 non-null object
17  open_acc         396030 non-null float64
18  pub_rec          396030 non-null float64
19  revol_bal        396030 non-null float64
20  revol_util       395754 non-null float64
21  total_acc        396030 non-null float64
22  initial_list_status 396030 non-null object
23  application_type  396030 non-null object
24  mort_acc         358235 non-null float64
25  pub_rec_bankruptcies 395495 non-null float64
26  address          396030 non-null object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB

```

```
# Checking null values
```

```
data.isnull().sum()
```

```

loan_amnt      0
term           0
int_rate       0
installment    0
grade          0
sub_grade      0
emp_title      22927
emp_length     18301
home_ownership 0
annual_inc     0
verification_status 0
issue_d        0
loan_status    0
purpose        0
title          1755
dti            0
earliest_cr_line 0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     276
total_acc      0
initial_list_status 0
application_type 0
mort_acc       37795
pub_rec_bankruptcies 535
address        0
dtype: int64

```

```
# Data cleaning
```

```
#----> We'll remove columns with only one unique value because their variance will be 0 and they won't help us anticipate anything.
```

```

features_nuniques = {}
for i in data.columns:
    features_nuniques[i] = data[i].nunique()

```

```
print(features_nuniques)
```

```
{'loan_amnt': 1397, 'term': 2, 'int_rate': 566, 'installment': 55706, 'grade': 7, 'sub_grade': 35, 'emp_title': 173105, 'emp_length': 18301, 'home_ownership': 5, 'annual_inc': 14730, 'verification_status': 4, 'issue_d': 1, 'loan_status': 10, 'purpose': 10, 'title': 1755, 'dti': 1, 'earliest_cr_line': 1, 'open_acc': 1, 'pub_rec': 1, 'revol_bal': 1, 'revol_util': 276, 'total_acc': 1, 'initial_list_status': 1, 'application_type': 1, 'mort_acc': 37795, 'pub_rec_bankruptcies': 535, 'address': 1}
```

```
## Since emp_title has null values more than 5%(5.7 %) and it does not influence the target variable, it is removed
```

```
data.drop('emp_title', axis = 1, inplace = True)
```

```

## Since only 0.06%, 0.4% , 0.1 & 4.6% of data are null values in revol_util,title, pub_rec_bankruptcies & employee_length respectively,
## Eventhough 9.5% of data are missed in mort_account,it is imputed with it's mean value as it is influence the target data
data['emp_length'].replace({"10+ years":"10 year", "< 1 year":"0 year", None:np.nan}, inplace = True)

```

```
data["emp_length"] = data["emp_length"].str.split(" ").str[0]
```

```
data["emp_length"] = data["emp_length"].astype('float64')
```

```

for i in ["emp_length","revol_util","mort_acc"]:
    data[i].fillna(data[i].mean(), inplace = True)

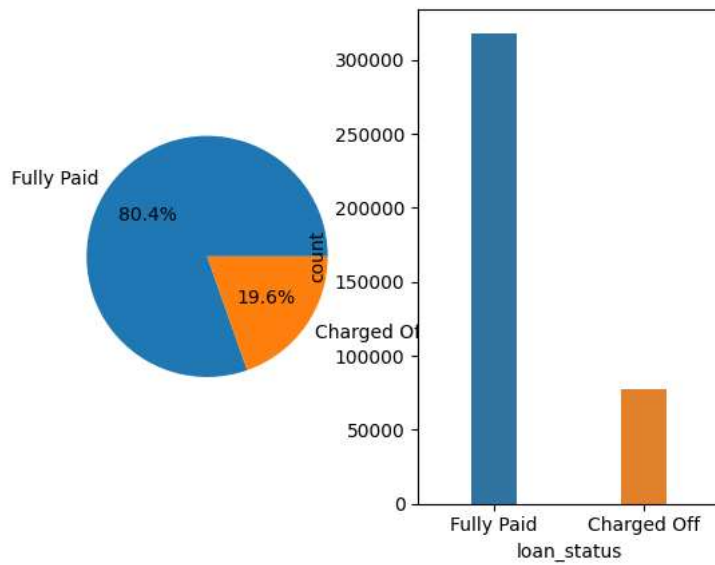
```

```
for j in ["title","pub_rec_bankruptcies"]:
```

```
    data[j].fillna(data[j].mode()[0], inplace = True)
```

```
# Exploratory data analysis
#--> target variable's(loan_status) distribution in the dataset
plt.subplot(1,2,1)
labels = ["Fully Paid", "Charged Off"]
plt.pie(data["loan_status"].value_counts(), labels = labels, autopct='%1.1f%%')
plt.subplot(1,2,2)
sns.countplot(data = data, x = "loan_status", width = 0.3)
```

<Axes: xlabel='loan_status', ylabel='count'>



```
# checking correlation between the variables
plt.figure(figsize=(15,8))
sns.heatmap(data.corr(), annot = True, linewidths = 2.5)
```

```
<ipython-input-99-490bb3ea7128>:3: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future ver
sns.heatmap(data.corr(), annot = True, linewidths = 2.5)
```

```
# Let's visualize if there is any relationship between the target variable and other variables
```

```
catedgarical_features = ["term", 'grade', 'sub_grade', 'home_ownership', 'verification_status', 'application_type']
```




```
for i in range(1, len(catedgarical_features)+1):
```

```
plt.subplot(3,2,i)
```


```
plt.xticks(rotation = 90)
```


```
sns.countplot(data = data, x = catedgarical_features[i-1], hue = "loan_status")
```

```
plt.show()
```



```
# Checkin impact of earliest_cr_line o target variable
data["earliest_cr_line"] = pd.to_datetime(data["earliest_cr_line"])
plt.figure(figsize = (15,15))
plt.xticks(rotation = 90)
sns.countplot(data = data, x = data["earliest_cr_line"].dt.year, hue = "loan_status")
```

 <Axes: xlabel='earliest_cr_line', ylabel='count'>



```
# Chisquare test- does earliest_cr_line really influence the loan status
year_loan_status = pd.crosstab(index = data["loan_status"], columns = data["earliest_cr_line"].dt.year)

year_loan_status
```

earliest_cr_line	1944	1948	1950	1951	1953	1954	1955	1956	1957	1958	...	20
loan_status												
Charged Off	1	0	0	1	0	0	3	3	2	4	...	45
Fully Paid	0	1	3	2	2	4	6	4	5	8	...	163

```

from scipy.stats import chi2_contingency

chi_stat, p_value, df, expected_freq = chi2_contingency(year_loan_status)
print(f"p_value is {p_value}")
if p_value < 0.05:
    print("earliest_cr_line has impact on loan status")
else:
    print("Has no impact")

    p_value is 6.1047661863873966e-198
    earliest_cr_line has impact on loan status

# Feature engineering
#--> Converting categorical target var into numerical

def cat_nume_target(x):
    if x == "Fully Paid":
        return 1
    else:
        return 0
data["loan_status"] = data["loan_status"].apply(cat_nume_target)

#--> Converting categorical features into numerical
# ----> term is nothing but the number of payments on the loan(36 months or 60 months)
data["term"] = data["term"].apply(lambda x: 36 if x == "36 months" else 60)

# -----> Since grade & subgrade have order, they are encoded using label encoder # verification type does not influence the target variable
from sklearn.preprocessing import OrdinalEncoder
cols = ["grade", "sub_grade", "initial_list_status"]
ordinal_encoder = OrdinalEncoder()
data[cols] = ordinal_encoder.fit_transform(data[cols])
# earliest_cr_line
data["earliest_cr_line"] = data["earliest_cr_line"].dt.year.astype(int)

#----> since home_ownership is not ordinal and only 5 categorical values it has, one-hot encoding is used to encode it

from sklearn.preprocessing import OneHotEncoder
col = ["home_ownership"]
ohe = OneHotEncoder(handle_unknown='ignore', sparse = False)

ohe_new = pd.DataFrame(ohe.fit_transform(data[col]))

# One-hot encoding removed index; put it back
ohe_new.index = data.index

# Remove categorical columns (will replace with one-hot encoding)
data = data.drop(col, axis=1)

# Add one-hot encoded columns to numerical features
data = pd.concat([data, ohe_new], axis=1)

# Ensure all columns have string type
data.columns = data.columns.astype(str)

/usr/local/lib/python3.10/dist-packages/sklearn/preprocessing/_encoders.py:868: FutureWarning: `sparse` was renamed to `sparse_output`
warnings.warn(

# Since verification_status, issue_d, purpose, title, application_type & address do not have influence on target variable, hence removed.
data.drop(["verification_status", "issue_d", "purpose", "title", "application_type", "address"], axis=1, inplace=True)

# train test split
x = data.drop("loan_status", axis=1)
y = data["loan_status"]

from sklearn.model_selection import train_test_split

x_train_cv, x_test, y_train_cv, y_test = train_test_split(x, y, test_size=0.2, random_state=1)

```

```
x_train, x_val, y_train, y_val = train_test_split(x_train_cv, y_train_cv, test_size = 0.2, random_state = 1)
```

```
x_train.shape
```

```
(253459, 24)
```

```
# Scaling all features in order to avoid dominance of any features due to its high range
```

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(x_train)
```

```
x_train = scaler.transform(x_train)
x_val = scaler.transform(x_val)
x_test = scaler.transform(x_test)
```

```
# Let's train a logistic regression model
```

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(x_train, y_train)
```

```
LogisticRegression
```

```
model.coef_
```

```
array([[ -5.07647543e-01,  0.00000000e+00,  2.34026322e-01,
         3.91174075e-01, -4.52658264e-04, -7.77469943e-01,
         1.18865997e-02,  2.04303398e-01, -5.33864276e-01,
         1.95553207e-02, -1.04545576e-01, -6.30875796e-02,
         7.12633515e-02, -8.18671391e-02,  1.00579255e-01,
        -1.04558597e-04,  7.29136053e-02,  2.58587365e-02,
         2.29695807e-02,  6.01427298e-02, -2.03460504e-03,
         4.87801806e-03, -7.41563208e-03, -5.71467588e-02]])
```

```
model.intercept_
```

```
array([1.57755842])
```

```
# Prediction on train data
model.predict(x_train)
```

```
array([1, 1, 1, ..., 1, 1, 1])
```

```
# Defining accuracy on train data
def accuracy(y_true, y_pred):
    return np.sum(y_true == y_pred) / (y_true.shape[0])
```

```
# Accuracy on train data
accuracy(y_train, model.predict(x_train))
```

```
0.8053491886261683
```

```
# Accuracy on test data
accuracy(y_test, model.predict(x_test))
```

```
0.8037269903795167
```

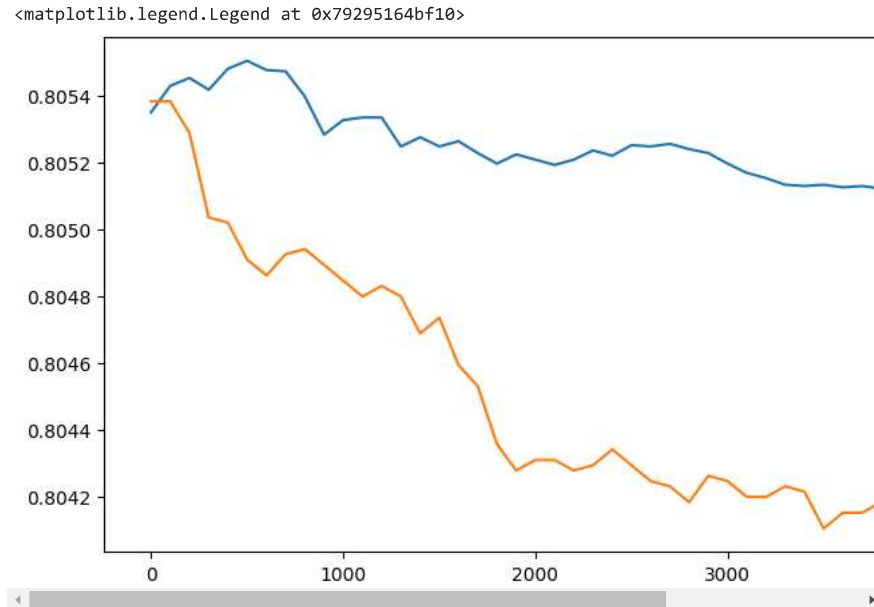
```
# Hyper parameter tuning (regularization strength)
```

```
from sklearn.pipeline import make_pipeline
```

```
train_scores = []
val_scores = []
```

```
scaler = StandardScaler()
for la in np.arange(0.01, 5000, 100):
    scaled_lr = make_pipeline(scaler, LogisticRegression(C = 1/la))
    scaled_lr.fit(x_train, y_train)
    train_score = accuracy(y_train, scaled_lr.predict(x_train))
    val_score = accuracy(y_val, scaled_lr.predict(x_val))
    train_scores.append(train_score)
    val_scores.append(val_score)
```

```
plt.figure(figsize = (10,5))
plt.plot(list(np.arange(0.01, 5000, 100)), train_scores, label= 'train')
plt.plot(list(np.arange(0.01, 5000, 100)), val_scores, label= 'val')
plt.legend(loc = 'lower right')
```



```
# As shown above, best regularization strength(lambda) is around 60
model2 = LogisticRegression(C = 1/70)
model2.fit(x_train,y_train)
```

```
LogisticRegression
LogisticRegression(C=0.014285714285714285)
```


```
accuracy(y_train, model2.predict(x_train))
```

```
0.8054399330858245
```

```
accuracy(y_test, model2.predict(x_test))
```

```
0.8037522409918441
```

```
# Evaluation of the performance of the model
#--> Confusionmatrix
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from matplotlib import pyplot as plt
y_pred = model2.predict(x_test)
conf_matrix = confusion_matrix(y_test,y_pred)
ConfusionMatrixDisplay(conf_matrix).plot()
plt.show()
```




```
from sklearn.metrics import recall_score, precision_score
```

```
print(f"recall is {recall_score(y_test, y_pred)}")
print("-----")
print(f"precision is {precision_score(y_test, y_pred)}")
print("-----")
print(f"accuracy is {accuracy(y_test, model2.predict(x_test))}")
```


```
recall is 0.9835735414437661
-----
precision is 0.8117176543915234
-----
accuracy is 0.8037522409918441
```

```
# Recall- precision tradeoff--- ROC curve
#---->For this particular case, precision is more important than recall
```

```
from sklearn.metrics import roc_curve, roc_auc_score, recall_score, precision_score
```

```
prob = model2.predict_proba(x_test)
proba1 = prob[:,1]
```

```
fpr, tpr, thresh = roc_curve(y_test, proba1)
plt.plot(fpr,tpr,"--", color = "red")
plt.title("ROC Curve")
plt.xlabel('FPR')
plt.ylabel("TPR")
plt.show()
```



```
# Area under the ROC curve
roc_auc_score(y_test, proba1)
```

```
0.7053091072912581
```

```
# Precision recall curve
```

```
from sklearn.metrics import precision_recall_curve
```

```
precision, recall, thresh = precision_recall_curve(y_test, proba1)
```

```
plt.plot(precision, recall,"--", color = "red")
plt.title("PR Curve")
plt.ylabel('precision')
plt.xlabel("recall")
plt.show()
```

